





Welcome to Week 12: Solving MDP

■ Outline:

- ▣ Solving MDP: Computing the best policy
 - ▣ The problem with Value iteration
 - ▣ Policy based method: Policy Evaluation + policy Improvement






Problems with Value Iteration





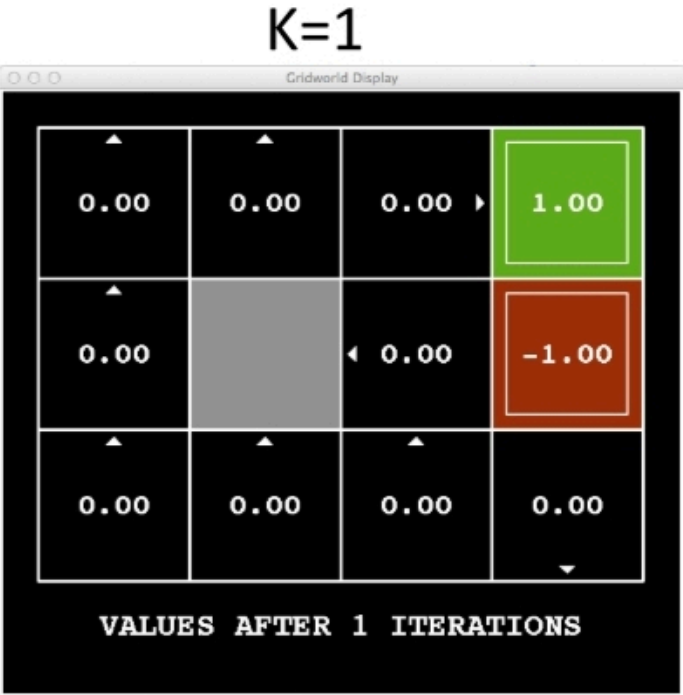
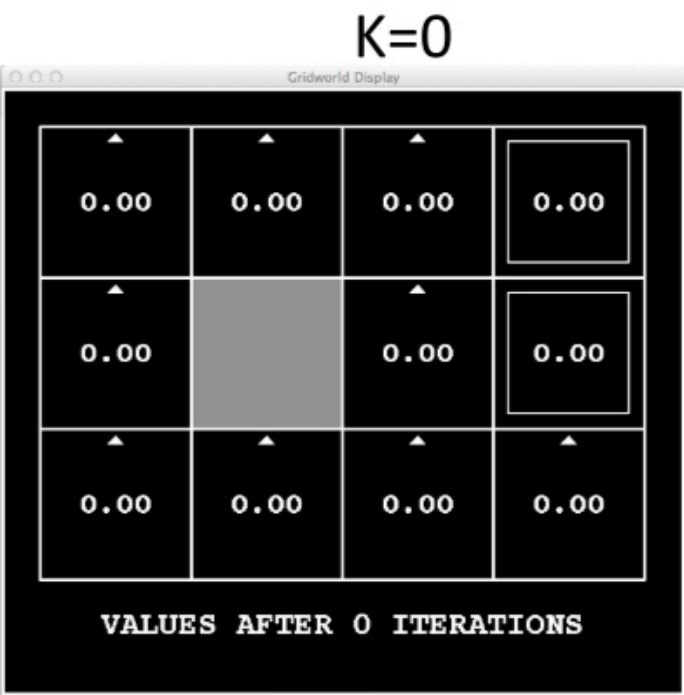
- Value iteration repeats the Bellman updates:

$$V_{k+1}(s) \leftarrow \max_a \sum_{s' \in \mathcal{S}} T(s, a, s') [R(s, a, s') + \gamma V_k^*(s')]$$

- **Problem 1:** It's slow. $O(|\mathcal{S}|^2 |A|)$ per iteration
- **Problem 2:** The max at each state rarely changes
The policy converges long before the values




The Policy Converges Long Before the Values



Policy Evaluation


■ In Bellman Equation, we didn't have a fixed policy.

□ Hence, we had to max over all actions in the Bellman equation.



$$V^*(s) = \max_{a \in A} \sum_{s' \in S} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

■ If we have a fixed policy π , what's the recursive relation of $V_\pi(s)$?



$$V_\pi(s) =$$

Policy Evaluation

■ How do we compute the V 's for a fixed policy π ?

□ **Idea 1:** Turn recursive Bellman equations into updates (like value iteration)

Initialize $V_{\pi,0}(s) = 0$, for all $s \in \mathcal{S}$

For $k=1$, until convergence:

For all $s \in \mathcal{S}$:

$$V_{\pi,k+1}(s) \leftarrow \sum_{s' \in \mathcal{S}} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_{\pi,k}(s')]$$

□ Computational Complexity per iteration:

□ **Idea 2:** With fixed policy, we just have a linear system of equations. Solve with your favorite linear solver.

Policy Iteration

- By Combining "policy evaluation" and "policy extraction", we design a nice policy based method to find optimal policy
- Policy Iteration:
 - Step 1: Policy Evaluation; Calculate values for some fixed policy (not optimal values!) until convergence.
 - Step 2: Policy Improvement; Given the values from step 1, update policy using one-step look-ahead
 - Repeat steps until convergence.

Policy Iteration

■ Policy Iteration:

Initialize $\pi_0(s)$ for all $s \in \mathcal{S}$

For $i=1$, until Convergence:

Policy Evaluation { Initialize $V_{\pi_i,0}(s)=0$ for all $s \in \mathcal{S}$

For $k=1$, until Convergence:

$$V_{\pi_i,k+1}(s) \leftarrow \sum_{s' \in \mathcal{S}} T(s, \pi_i(s), s') [R(s, \pi_i(s), s') + \gamma V_{\pi_i,k}(s')]$$

Store the values after convergence as $V_{\pi_i}(s)$.

Policy Improvement { For all $s \in \mathcal{S}$:

$$\pi_{i+1}(s) \leftarrow \arg \max_{a \in \mathcal{A}} \sum_{s' \in \mathcal{S}} T(s, a, s') [R(s, a, s') + \gamma V_{\pi_i}(s')]$$

Summary

- To Compute the optimal values of an MDP, use
 - ▣ Value iteration
 - ▣ or, Policy iteration
- To Compute values for a fixed policy, use
 - ▣ Policy evaluation
- To turn your values into a policy, use
 - ▣ Policy extraction

Summary: They are All Just Variation of The Same Equation

- Optimal V and Q functions:

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$(V^*(s) = \max_a Q^*(s, a))$$

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma \max_{a'} Q^*(s', a')]$$

- Value function for a fixed policy:

$$V_\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_\pi(s')]$$

- Policy π^* for V^* and Q^* function:

$$\pi^*(s) = \arg \max_a Q^*(s, a)$$

$$\pi^*(s) = \arg \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

Comparison

- Both value iteration and Policy iteration Compute the same thing
 - In value iteration
 - Every iteration updates both the values and (implicitly) the policy
 - We don't track the policy, but taking the max over actions implicitly recomputes it
 - In policy iteration:
 - We do several passes that update utilities with fixed policy
 - After the policy is evaluated, a new policy is chosen
 - The new policy will be better (or we're done)
- Both methods are dynamic programs to solve MDP's
- each pass is fast because no max.

