

# Welcome to Week 12: Solving MDP

## ■ Outline:


- ▣ Review: Computing value of states, given policy  $\pi$
- ▣ Dynamic Programming: An Iterative Algorithm to find the values given a policy
- ▣ Solving MDP: Computing the best policy
  - ▣ Value iteration and policy Extraction
  - ▣

■ ▣ ▣

■ ▣ ▣

■ ▣ ▣

# Review: Markov Property Provides a Nice Structure






with probability  $T(s, \pi(s), s')$ , you will receive the immediate reward  $R(s, \pi(s), s')$  and go to the next state  $s'$ . From state  $s'$ , your expected utility that you would receive is  $V_\pi(s')$ . This is the reward you receive in the next step. Hence, you should discount it by  $\gamma$

$$\begin{aligned}
 V_\pi(s) = & T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_\pi(s')] \\
 & + T(s, \pi(s), s'') [R(s, \pi(s), s'') + \gamma V_\pi(s'')] \\
 & + T(s, \pi(s), s''') [R(s, \pi(s), s''') + \gamma V_\pi(s''')]
 \end{aligned}$$

# Markov Property Provides a Nice Structure

- Given Policy  $\pi$ , the value of states satisfy

$$V_{\pi}(s) = \sum_{s' \in S} \underbrace{T(s, \pi(s), s')}_{\text{probability of this transition}} \left[ \underbrace{R(s, \pi(s), s')}_{\text{immediate Reward for this transition}} + \underbrace{\gamma V_{\pi}(s')}_{\text{Discounted Expected Return after this transition}} \right], \quad \forall s \in S$$



# Markov Property Provides a Nice Structure

- For finite state MDP, we can express the previous equations as a matrix equation

$$\text{Let } \underline{v} = \begin{bmatrix} v_\pi(s) \\ v_\pi(s') \\ \vdots \end{bmatrix}, \quad T = \begin{bmatrix} T(s, \pi(s), s) & T(s, \pi(s), s') & \dots \\ T(s', \pi(s'), s) & T(s', \pi(s'), s') & \dots \\ \vdots & \ddots & \ddots \end{bmatrix}, \quad R = \begin{bmatrix} R(s, \pi(s), s) & R(s, \pi(s), s') & \dots \\ R(s', \pi(s'), s) & R(s', \pi(s'), s') & \dots \\ \vdots & \ddots & \ddots \end{bmatrix}$$

$$\underline{v} = (T \odot R) \underline{1} + \gamma T \underline{v}$$

- Thus,  $\underline{v} = (I - \gamma T)^{-1} (T \odot R) \underline{1}$
- Solving directly requires taking a matrix inverse  $\sim O(|S|^3)$

# Iterative Algorithm for Computing Values, Given a Policy

## ■ Dynamic Programming:

Initialize  $V_{\pi,0}(s) = 0$  for all  $s \in \mathcal{S}$

For  $k=1$ , until Convergence:




For all  $s \in \mathcal{S}$ :

$$V_{\pi,k}(s) = \sum_{s' \in \mathcal{S}} T(s, \pi(s), s') \left[ R(s, \pi(s), s') + \gamma V_{\pi,k-1}(s') \right]$$

## ■ Computational Complexity For Each Iteration:


# Solving MDP

- An MDP can be specified with a tuple:  $(S, A, T, R, \gamma)$
- To solve an MDP means:
- What is the definition of the "best" policy?



# Example: Solving Mars Rover in 2D-Grid

- Consider the following Mars Rover in 2D Grid.



Mars Rover gets a reward of (+1) if it gets to this location. Then, game ends.



Mars Rover gets a reward of (-1) if it gets to this location. Then, game ends.

□ Noise = 0.2 , Discount = 0.9 , Living reward = 0


- What would the solution of this MDP look like?





# 2D-Grid Mars Rover: Optimal Policy

- Optimal Policy  $\pi^*$ :
- $V$  values of the optimal policy:



- Q-value of state-action pair:  
The expected utility of taking action "a" from state "s", and then following the policy







-  
-  







# Values of States: Bellman Equation



■ Recursive definition of values:



■  



■  



■  

■  


■  

■  

■  

■  

■ But how do we solve these equations?



# Value Iteration Algorithm

- Value iteration:

Initialize  $V_0(s) = 0$  for all  $s \in S$

For  $k=1$ , until Convergence:

For  $s \in S$ :








$$V_{k+1} \leftarrow \max_{a \in A} \sum_{s' \in S} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

- Applies the Bellman Equation to update the value of state  $s$  based on the state values derived in the previous iteration

- Repeat until convergence, which yields  $V^*$

- Complexity of each iteration:

# Example: Mars Rover in the 2D-Grid



# Homework: Run Value Iteration on "Erfan the Undergrad" Example

- After how many iterations of "Erfan the Undergrad" you will reach the optimal values?


# Bellman Equation for Q-Values

- We saw the Bellman Equation for optimal  $V(s)^*$

$$V(s)^* = \max_{a \in A} \sum_{s' \in S} T(s, a, s') [R(s, a, s') + \gamma V(s')^*]$$

- Can you write down Bellman Equation for  $Q^*(s, a)$ ?

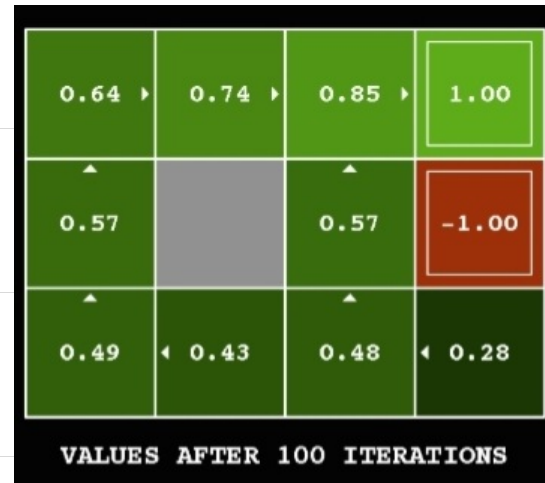
$$Q^*(s, a) =$$




- Leads to Q-value iteration we will see later

# Policy Extraction

- So far, with value iteration, we could find the optimal values




- How can we use the optimal values to find the optimal policy?



# Computing Actions from Values

- The action  $a$  that maximizes  $Q^*(s, a)$  is the optimal action  $\pi^*(s) =$



- This process is called **policy extraction**.
- If we had access to q-values, we could simply find  $\pi^*(s)$  as
- **Important lesson:** Actions are easier to select from q-values than values!

# Welcome to Week 12: Solving MDP

