

Active Learning

- We saw how to learn $V_{\pi}(s)$ with TD Learning.
- That was good for passive Learning.
 - In passive Learning, we are given samples generated by policy π
 - Our task is to find $V_{\pi}(s)$. TD Learning is a model-free approach to learn $V_{\pi}(s)$'s.
- **Active Learning:** We want to be able to do more than that?
We want to be able to update policy π (i.e., the sampling policy) so that not only we collect samples and learn the model, but also make decision that give us good rewards.

Active Learning

■ Active Reinforcement Learning:

☐ Actively collecting data, while we learn the model

■ Problem model:

☐ We don't know T and R , and we choose the actions.

☐ Goal: Learn the optimal policy / actions.

Active Learning

- In active reinforcement learning, **Learner** makes the choices.
- **Fundamental tradeoff:**
 - exploration v.s. exploitation

Recall: To Estimate an Expectation, use Running Average of Samples

■ Let's first write the Bellman equation in terms of Values (i.e. $V^*(s)$):

$$\boxed{\square} V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$\boxed{\square}$ Value iteration finds $V^*(s)$ with dynamic programming, i.e.

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

$\boxed{\square}$ In reinforcement learning, we don't have T and R .

★ That's why we usually use samples to estimate expectations.

● e.g., in TD we could estimate $V_{\pi}(s)$ w/ running average of samples

$$V_{\pi}(s) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_{\pi}(s')] \xrightarrow[\text{running average of samples}]{\text{So, we estimated it with}} \bar{V}_{\pi, k+1}(s) = (1-\alpha) \bar{V}_{\pi, k}(s) + \alpha [R(s, a, s') + \gamma \bar{V}_{\pi, k}(s')]$$

Can We Use Running Average to Estimate Optimal Values?

■ Now, we want to do active learning,

□ i.e., learning the optimal policy / value / q-value

■ Can we use the same idea as in TD to find $V^*(s)$?

□ In other words, can we use running averaging to find $V^*(s)$?

■ Let's revisit the recursion for optimal values, i.e. $V^*(s)$:

$$\square V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

■ Main question: Is $V^*(s)$ written in the above recursion an expectation? Yes ☐ No ☒


why? Because of the max.

Can We Use Running Average to Estimate q-values?

■ Now, let's write the recursion for $Q^*(s,a)$ in terms of Q^*

$$\square Q^*(s,a) = \sum_{s'} T(s,a,s') \left[R(s,a,s') + \gamma \max_{a'} Q^*(s',a') \right]$$

This is an expectation



■ Nice! We can use the running average:

$$\text{Sample} = R(s,a,s') + \gamma \max_{a'} Q(s',a')$$

$$Q(s,a) \leftarrow (1-\alpha) Q(s,a) + \alpha \text{ Sample}$$

Q-learning

■ Q-learning: Sample-based q-value iteration.

■ We have: $Q^*(s,a) = \sum T(s,a,s') [R(s,a,s') + \gamma \max_{a'} Q^*(s',a')]$

■ Thus, we can estimate the above expectation by

□ Whenever you receive a new sample from a transition like $(s, a, s', R(s,a,s'))$,

□ By using running averaging you can update your estimate

of $Q^*(s,a)$: $Q_{k+1}(s,a) \leftarrow (1-\alpha_t) Q_k(s,a) + \alpha_t \underbrace{[R(s,a,s') + \gamma \max_{a'} Q_k(s',a')]}_{\text{new sample}}$

Q-Learning Properties

- Q-learning Converges to optimal q -values, even if you're acting suboptimally!
- This is called **off-policy Learning**.
- what it means is that in the limit it doesn't matter how you select the action. Regardless of action selection, you will converge to the optimal q -values (hence optimal policy) in the limit.
- But,
 - You have to explore enough (Needs to visit all states often enough)
 - you have to eventually make the learning rate small enough, but not decrease too quickly. $\sum \alpha_t = \infty$, $\sum \alpha_t^2 < \infty$

Exploration vs. Exploitation

- It's not a good idea to always **exploit** the action that you consider to be the best action you've seen so far.
- You have to try many actions. You must **explore**.
 - They may not be good for you. But you won't know unless you try them.

How to Explore

■ Several schemes for forcing exploration.

□ random actions (ϵ -greedy) $0 < \epsilon \leq 1$

- Every step, flip a coin.

- with (small) ϵ probability, act randomly

- with $(1-\epsilon)$ probability, act based on the current policy (i.e. estimate q -values)

□ with ϵ -greedy, we eventually explore the space \rightarrow hence, find the optimal values

But, it keeps randomly choosing actions even when the learning is done

□ Solutions: i) lower ϵ over time ii) use exploration function

i) Lower ϵ Over Time

- One option is to set $\epsilon = 1/t$.
- i.e., at time step t , with probability $1/t$, act randomly
- with probability $(1 - 1/t)$, act on optimal policy.

■ It does eventually converge, but

↳ It can be slow

↳ It explore all states with similar probability

ii) Exploration Functions

■ Idea: to explore areas whose badness is not (yet) established, and eventually stop exploring

■ We can implement this idea by modifying our estimate of the q -values.

□ increase the estimate q -values based on how unexplored they are:

- previously, new sample was: $R(s, a, s') + \gamma \max_{a'} Q(s'; a')$
- with exploration function, we modify the new sample to $R(s, a, s') + \gamma \max_{a'} f(Q(s'; a'), N(s'; a'))$

Exploration Functions, cont'd

■ $N(s', a')$: # times we had seen the pair (s', a') before

■ $f(Q(s', a'), N(s', a')) = Q(s', a') + \frac{K}{N(s', a')}$ ← Some Constant

↪ exploration function

■ So, the update rule with exploration function will be

$$Q_{k+1}(s, a) \leftarrow (1 - \alpha) Q_k(s, a) + \alpha \left[R(s, a, s') + \gamma \max_{a'} f(Q(s', a'), N(s', a')) \right]$$

■ **Note:** This propagates the "bonus" back to states that lead to unknown states, as well.